

BigFive5q Predictive Test

Juan Eloy Suarez - PH125 Data Science - Harvard

24/June/2021

Executive Summary

This project uses prediction techniques to simplify the well-known “Big Five” personality traits determination test. The goal is to predict most of the answers from a person so we can reduce dramatically his/her input whilst getting to a final score as accurate as possible. Thus, less length, effort and time of test resolution will make it easier to be integrated in other processes where estimation of personality characteristics of an individual can be useful.

One of the key aspects of the project is to define clear and meaningful *indicators of accuracy* that take into account complexity to measure the multi-outcome challenge. Related with this, a Montecarlo approach has been used to simulate a base-reference of random results.

Project originality roots on the fact of taking advantage of a well-known standard **recommendation model**, adapting some concepts to use classical “User x Item” as inspiration for a “User x Question/Trait” structure.

Apart of de Exploratory Data Analysis and training of a operational prediction model, an **interactive web app** (<https://juanesuarez.shinyapps.io/BigFive5q>) has been created for illustrative purpose, where user can *play* with the answers and even interactively see the hidden predictions obtained for unanswered questions.

1 Introduction

The Big Five personality traits, also known as the five-factor model (FFM) and the OCEAN model, is a taxonomy, or grouping, for personality traits. When factor analysis (a statistical technique) is applied to personality survey data, some words used to describe aspects of personality are often applied to the same person. For example, someone described as conscientious is more likely to be described as “always prepared” rather than “messy”.

This theory is based therefore on the association between words but not on neuropsychological experiments, by using descriptors of common language. As a result, it suggests five broad dimensions commonly used to describe the human personality and psyche.

This test is made of 50 questions grouped in five *personality traits*. All questions must be answered from 1 (Strongly Disagree) to 5 (Strongly Agree), being some of them “reverted” (formulated as negative). The percentile of the average of each 10-questions group is the final score for that trait. Our **goal is to use only five questions to predict all 45 others and therefore final scores for each trait**, benefiting therefore of such a very light approach to get results.

Trait	Associated Tendencies	
Extraversion (vs. <i>Introversion</i>)	· Gregarious (sociable) · Assertiveness (forceful) · Activity (energetic)	· Excitement-Seeking (adventurous) · Positive Emotions (enthusiastic) · Warmth (outgoing)
Conscientiousness (vs. <i>Lack of Direction</i>)	· Competence (efficient) · Order (organized) · Dutifulness (not careless)	· Achievement striving (thorough) · Self-disciplined (not lazy) · Deliberation (not impulsive)
Openness to Experience (vs. <i>Closedness</i>)	· Ideas (curious) · Fantasy (imaginative) · Aesthetics (artistic)	· Actions (wide interests) · Feelings (excitable) · Values (unconventional)
Agreeableness (vs. <i>Antagonism</i>)	· Trust (forgiving) · Modesty (not show-off) · Altruism (warm)	· Compliance (not stubborn) · Tender-mindedness (sympathetic) · Straightforwardness (not demanding)
Neuroticism (vs. <i>Emotional Stability</i>)	· Anxiety (tense) · Angry Hostility (irritable) · Depression (not contented)	· Self-consciousness (shy) · Impulsiveness (moody) · Vulnerability (not self-confident)

Source: Adapted from Costa and McCrae, 1992

Figure 1: The Big Five Framework of Personality Traits

Source code and documentation

All project contents, source codes, documentation, data and shinnyapp codes are publicly available at: <https://github.com/JuanESuarez/BigFive5q/tree/main>

2 Exploratory Data Analysis

The “Big Five Personality Test. 1M Answers to 50 personality items, and technical information” dataset in our source for this project. The personality test was constructed with the “Big-Five Factor Markers” from the International Personality Item Pool (<https://ipip.ori.org/newBigFive5broadKey.htm>). Participants were informed that their responses would be recorded and used for research at the beginning of the test, and asked to confirm their consent at the end of the test. These data were collected (2016-2018) through an interactive on-line personality test and it contains:

- 1015341 records (tests).
- 110 columns of data, 50 of them answers (1=Strongly Disagree, 2=Disagree, 3=Neutral, 4=Agree, 5=Strongly Agree) of test questions, and all others containing country, location, load date and load/response time for each question.
- A codebook detailing literal wording of each one of the 50 questions.

For our approach to the model we will discard country, location, and load/response time to questions, so we focus directly on the questions and their answers. As we will remark in the Conclusions chapter, there are interesting possibilities for developing the project using these data, but they are out of the scope of this work. This reduction in the number of prescriptors will significantly reduce the size of our dataset.

Data load process

Data are located in Kaggle site (<https://www.kaggle.com/tunguz/big-five-personality-test/download>). Data size in .CSV is 396 MB, so we will, as a previous step, load, select and convert them to the much more efficient format .rds, obtaining a very significant size reduction is obtained to 21 MB full dataset. The code to connect directly to Kaggle server requires an API and some specific coding, so for simplification, we have previously downloaded data locally and saved to a .rds file in directory ./data-source. Code to perform this preprocess and storage is available and shpwn below:

```
# Read Kaggle dataset copied to local directory Answers to
# test data tab delimited, and our int columns are
# character, perhaps because of the nulls!
df <- read.csv("./data_source/data-final.csv", sep = "\t", stringsAsFactors = FALSE,
  na.strings = c("NA", "NaN", " ", "NULL"))
# We select relevant columns to use
df <- df %>%
  select(c(1:50), dateload, country) %>%
  rownames_to_column("userId") %>%
  mutate(Month = month(dateload), Year = year(dateload)) %>%
  select(-dateload)
# Remove rows containing any NA
df <- na.omit(df)
# Remove any value not between 1 and 5 in the answers
# columns
df <- df %>%
  filter_at(vars(2:51), all_vars((.) %in% c(1:5)))
# Let's also load up the questions from the data dictionary
dictionary <- read_table("./data_source/codebook.txt", skip = 5) %>%
  separate(1, sep = "\t", extra = "merge", into = c("ID", "Question")) %>%
  data.frame() %>%
  top_n(50)
# Save .RDS input data to local files
```

```
saveRDS(df, "./data_source/BFtests.rds")
saveRDS(dictionary, "./data_source/dictionary.rds")
```

However, it is not necessary to run the this preload to run project, since all two referred input datasets are already, and publicly, available:

```
./data-source/BFdata.rds
./data-source/dictionary.rds
```

Distribution of scores per trait

We will refer as *Scores* to the result (in percentile) of the average of answers (1 to 5 as possible values) corresponding to the ten questions forming each trait. This must be done considering that part of the questions are formulated “reverted” according to published test definition.

```
# This list identifies each column to its corresponding
# group and informs of the sign (to revert if necessary)
# based in our previous analysis of correlation
keys.list <- list(openness = c("OPN1", "-OPN2", "OPN3", "-OPN4",
  "OPN5", "-OPN6", "OPN7", "OPN8", "OPN9", "OPN10"), conscienciousness = c("CSN1",
  "-CSN2", "CSN3", "-CSN4", "CSN5", "-CSN6", "CSN7", "-CSN8",
  "CSN9", "CSN10"), extroversion = c("EXT1", "-EXT2", "EXT3",
  "-EXT4", "EXT5", "-EXT6", "EXT7", "-EXT8", "EXT9", "-EXT10"),
  agreeability = c("-AGR1", "AGR2", "-AGR3", "AGR4", "-AGR5",
  "AGR6", "-AGR7", "AGR8", "AGR9", "AGR10"), natural_reactions = c("EST1",
  "-EST2", "EST3", "-EST4", "EST5", "EST6", "EST7", "EST8",
  "EST9", "EST10"))
```

Since this type of scoring is very frequent in psychometrics, we will take advantage of the package **psych**. Its function **scoreItems** will directly calculate scores (i.e. mean per group taking sign into account). Detailed documentation on this function can be found: See: <https://www.rdocumentation.org/packages/psych/versions/2.1.3/topics/scoreItems>

```
scoresfile <- scoreFast(keys.list, df) #scoresfile <- scoreVeryFast(keys.list, df)
# Add calculated columns (scores)
df_scored_preReversion <- cbind(df, scoresfile)
# Once obtained the observation average per group, we need
# the p-value (0 to 100) per observation in group scope
P1 = ecdf(scoresfile[,1]) # P is a function giving the empirical CDF of X
P2 = ecdf(scoresfile[,2])
P3 = ecdf(scoresfile[,3])
P4 = ecdf(scoresfile[,4])
P5 = ecdf(scoresfile[,5])
# Store a rounded percentage (percentile) as resulting score for each trait
df_scored_preReversion <- df_scored_preReversion %>%
  mutate(
    openness-P = round(P1('openness-A')*100),
    conscienciousness-P = round(P2('conscienciousness-A')*100),
    extroversion-P = round(P3('extroversion-A')*100),
    agreeability-P = round(P4('agreeability-A')*100),
    natural_reactions-P = round(P5('natural_reactions-A')*100)
  )
```

```
## Error: <text>:14:15: inesperado '='
## 13:   mutate(
## 14:     openness-P =
##           ^
```

We now visualize how some of the traits tend to be valued higher than others. For a global analysis, a normalization should be adequate, but we are now focusing in individual study of questions & traits, so the most significant interpretation is that each one of them is approximately normal, which be useful later for percentile calculations.

```
## Error in eval(lhs, parent, parent): objeto 'df_scored_preReversion' no encontrado
```

Test Reliability in each personality trait (internal consistency)

We are now interested in assessing the internal consistency of the test, i.e. to know **at what extent our questions are related, behave similarly, when they belong to same group/trait**.

The Cronbach's alpha coefficient measures reliability, or internal consistency, to see if multiple-question Likert scale surveys are reliable. Cronbach's alpha will tell us how closely related a set of test items are as a group. More details on Cronbach's alpha can be found here: <https://www.rdocumentation.org/packages/psych/versions/2.1.3/topics/alpha>

According to this coefficient generalized interpretation, values of alpha between 0.8 and 0.9 are considered "Good Internal Consistency", while alpha greater than 0.9 means an "Excellent Internal consistency" situation.

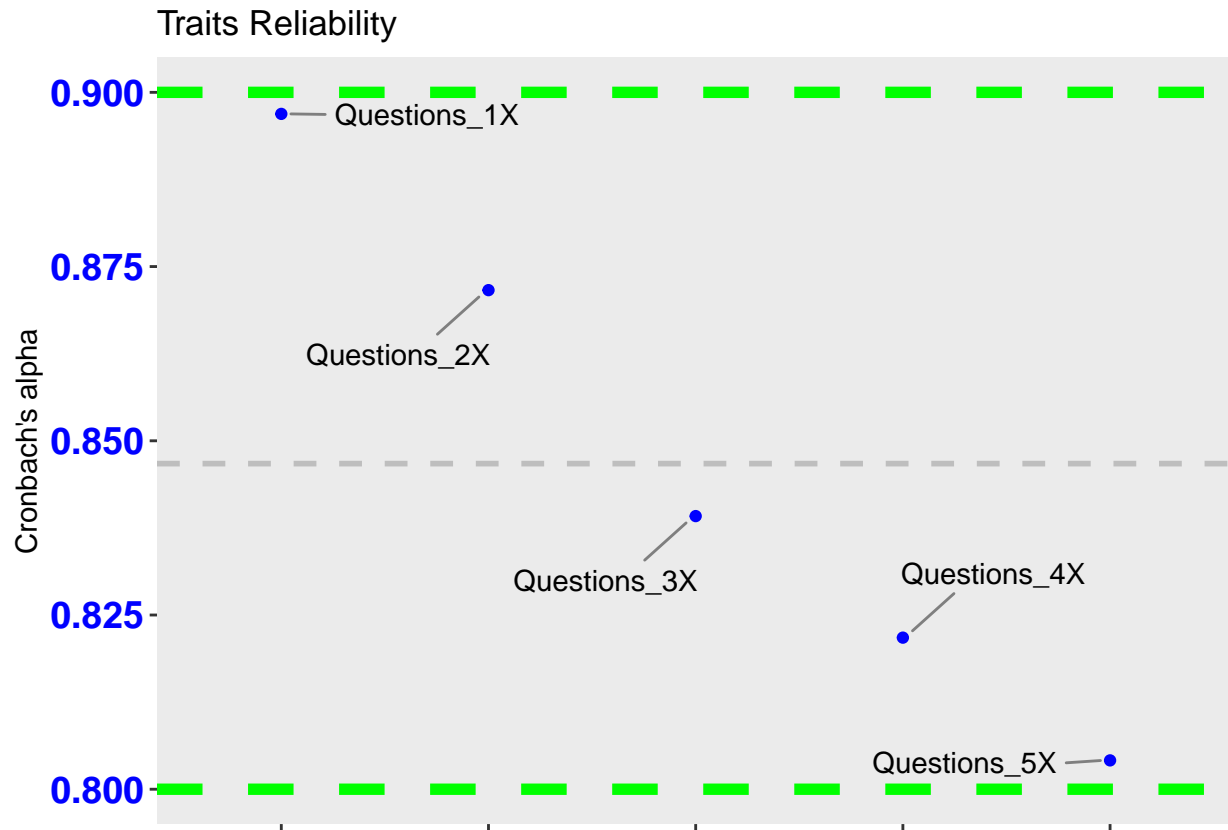
```
# Preparation, questions structure To prepare an analysis
# per 'traits' (each one includes groups of questions), we
# create a list of 5 concepts (groups of questions per
# personality dimension)
extroversion = dictionary[c(1:10), "ID"]
natural_reactions = dictionary[c(11:20), "ID"]
agreeability = dictionary[c(21:30), "ID"]
conscientiousness = dictionary[c(31:40), "ID"]
openness = dictionary[c(41:50), "ID"]
buckets = list(extroversion = extroversion, natural_reactions = natural_reactions,
               agreeability = agreeability, conscientiousness = conscientiousness,
               openness = openness)

# Groups analysis For each concept, let's calculate alfa to
# check 'internal consistency' of its questions The result
# also identify which variables should be inverted
# (negative correlation)
questionsSigns <- numeric()
traitsAlphas <- data.frame()
for (i in 1:length(buckets)) {
  questiondf <- df[, colnames(df) %in% unlist(buckets[i])]
  cronbach <- psych::alpha(questiondf, check.keys = TRUE)
  questionsSigns <- c(questionsSigns, cronbach$keys)
  traitsAlphas <- rbind(traitsAlphas, cronbach$total[, 1:3])
}
```

We also see extracted signs fro each question indicating if correlates directly or "reverted"

##	EXT1	EXT2	EXT3	EXT4	EXT5	EXT6	EXT7	EXT8	EXT9	EXT10	EST1	EST2	EST3
##	1	-1	1	-1	1	-1	1	-1	1	-1	1	-1	1
##	EST4	EST5	EST6	EST7	EST8	EST9	EST10	AGR1	AGR2	AGR3	AGR4	AGR5	AGR6
##	-1	1	1	1	1	1	1	-1	1	-1	1	-1	1
##	AGR7	AGR8	AGR9	AGR10	CSN1	CSN2	CSN3	CSN4	CSN5	CSN6	CSN7	CSN8	CSN9
##	-1	1	1	1	1	-1	1	-1	1	-1	1	-1	1
##	CSN10	OPN1	OPN2	OPN3	OPN4	OPN5	OPN6	OPN7	OPN8	OPN9	OPN10		
##	1	1	-1	1	-1	1	-1	1	1	1	1		

We easily visualize that, despite of some differences among traits, all of them fall in an, at least, “high integrity” interval:



Correlation of the answers to questions

A key aspect for understanding how some questions relate to other, even belong to a different group, is their correlation. It is intuitive to think in **correlation of questions as a clear indicator of how unknown values (unanswered question) can be predicted** based in known answers to other questions that usually correlate.

So, let's calculate correlation matrix among all questions to visualize global data pattern. Thus, we have also a have a good source for further calculations on what questions must be chosen to answers when there is a possibility to do it. Note some questions are reverted, but this is not relevant for this exploratory analysis, since it affects only to the sign of the correlation, not its absolute value.

A sample of the correlation matrix for five questions:

```

# Correlation calculation. Calculate global correlation
# matrix to see full landscape instead of just per trait
corAllQuestions <- cor(df[, 2:51])
# Sample of values (two per trait)
head(corAllQuestions[, c(1, 11, 21, 31, 41)])

```

```

##          EXT1      EST1      AGR1      CSN1      OPN1
## EXT1  1.00000 -0.101484 -0.027445  0.0261155  0.022487
## EXT2 -0.46953  0.050107  0.122157 -0.0016617 -0.052409
## EXT3  0.47293 -0.230392 -0.115756  0.1165689  0.014448
## EXT4 -0.52578  0.143897  0.084336 -0.0364316 -0.036729
## EXT5  0.49418 -0.121487 -0.107005  0.0859311  0.073811
## EXT6 -0.34675  0.091151  0.162322 -0.0465385 -0.190383


```


```

# Significance calculation. This function calculates de
# significance test associated to the correlation matrix
cor.mtest <- function(mat, ...) {
  mat <- as.matrix(mat)
  n <- ncol(mat)
  p.mat <- matrix(NA, n, n)
  diag(p.mat) <- 0
  for (i in 1:(n - 1)) {
    for (j in (i + 1):n) {
      tmp <- cor.test(mat[, i], mat[, j], ...)
      p.mat[i, j] <- p.mat[j, i] <- tmp$p.value
    }
  }
  colnames(p.mat) <- rownames(p.mat) <- colnames(mat)
  p.mat
}
p.mat <- cor.mtest(corAllQuestions) # significance test

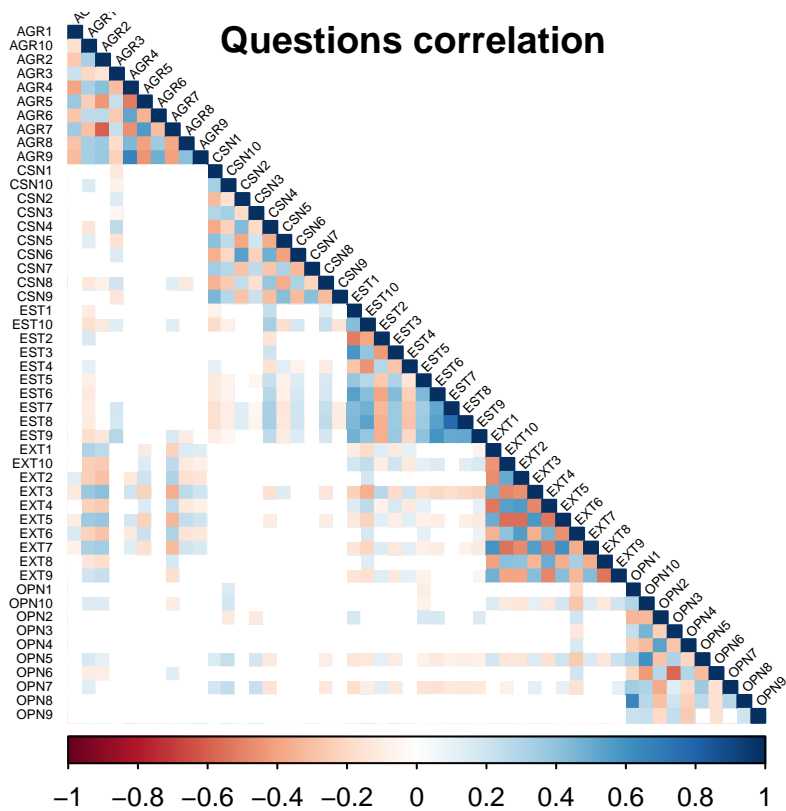
```

Correlation within same group: Questions explicitly associated to each “trait” strongly correlate, i.e. all of them are related because all of them explain the final value of the score for its specific group (trait). We see different traits (groups). Correlation is clearly high within same group.





Regarding behaviour among groups, we see also some other **significant correlations** (and of course also lack of them) of some questions with questions that “belong” to different traits. Since our challenge is precisely to use few (only five) questions to explain as much as possible of the result for all traits, it will be useful to **use these correlations to select what specific questions we show to get answer.**



3 Methods and train

3.1 Strategy, process

Our challenge is to predict the answers of 45 questions based on the answers to 5 questions we know. Total number of questions is exactly 50, so our matrix of users versus item (questions) is fully filled and we find no *sparsity*. We also observe that the shape of this matrix is pretty vertical, with around rows by only fifty columns. Despite of this singularities, this model seems to reasonable fit a “recommendation system” approach. We will pursue on that and take advantage of the already prebould strong tools existing in R, specifically the *Recommenderlab* package.

Nevertheless, some preparation steps will be necessary to *adapt* our problem to the functionalities the standard recommender tool provides.

Firstly, we need to prepare data for modeling, taking into account that our recommender algorithm will ignore **reverted effect on questions**, so we need to invert their existing answers in advance to make data homogeneous. Also, the **size of base dataset** is excessive in terms of computability in a local environment so we will make a random reduction of it - due to the high density of our Users x Questions matrix, this size reduction will not have dramatics effect in the decrease of our results, though of course could be reconfigured if more powerful systems would run the model.

We will also, in support of our interactive app, to extract data to make it **available within de *shinyapps* interactive web app**.

A key step will be then to separate dataset in two (**train and validation**) partitions for fulfill machine learning paradigm. For clarity, we will not use *scheme* options of recommenderlab package but to directly build, manage and measure these partitions.

Once data are ready, training, prediction and measurement process starts. Initially, we need to clearly define **what is and how to measure accuracy** of our model, as easily as possible, taking into account the relative complexity of the calculation of the five resulting scores. A very important step of the modeling will be to build an **algorithm to select what five questions we want to retrieve** answers from. This process must respect a random base to ensure eventual variability in results capture, but optimizing the choice to improve model accuracy. Once done all above steps, we will **run our recommender algorithms**, doing training of each model, prediction, calculation of final scores and storing results for comparison.

3.2 Data preparation for modeling

Reverted questions treatment

Our data contains answers to questions considered “reverted”. This means that the question is written in a negative way and, for global analysis, must be scored reverting results recorded. For example, questions “I am the life of the party” and “I don’t talk a lot” belong to same trait “Extraversion”. However, first second one is reverted, so for direct comparison we need to convert 1 to 5, 2 to 4, 4 to 2 and 5 to 1 answers.

This conversion is already covered by the generic scoring analysis used for package PSYCH, and just identifying with a minus sign before question Id, function internally reverts answers value. This was used in the EDA part above.

However, we are going to use functions based in recommendation methods (recommenderLab) based in linear algebra and distances. This algorithms use “ratings” always meaning “positive”, never “reverted”. For this reason, we will, before analysis, revert results of “reverse” questions and when necessary keep using regular PSYCH functions, but marking those questions as regular/positive (removing the minus sign).

```

# Revert negative questions of a Likert Scale: (n+1)-x
df <- df %>%
  mutate(OPN2 = abs(6 - OPN2), OPN4 = abs(6 - OPN4), OPN6 = abs(6 -
    OPN6), CSN2 = abs(6 - CSN2), CSN4 = abs(6 - CSN4), CSN6 = abs(6 -
    CSN6), CSN8 = abs(6 - CSN8), EXT2 = abs(6 - EXT2), EXT4 = abs(6 -
    EXT4), EXT6 = abs(6 - EXT6), EXT8 = abs(6 - EXT8), EXT10 = abs(6 -
    EXT10), AGR1 = abs(6 - AGR1), AGR3 = abs(6 - AGR3), AGR5 = abs(6 -
    AGR5), AGR7 = abs(6 - AGR7), EST2 = abs(6 - EST2), EST4 = abs(6 -
    EST4))
# We score AFTER change of 'reverted questions' Removing
# minus sign of the list. All questions will be 'positive'
keys.list.allPositive <- list(openess = c("OPN1", "OPN2", "OPN3",
  "OPN4", "OPN5", "OPN6", "OPN7", "OPN8", "OPN9", "OPN10"),
  conscienciousness = c("CSN1", "CSN2", "CSN3", "CSN4", "CSN5",
    "CSN6", "CSN7", "CSN8", "CSN9", "CSN10"), extroversion = c("EXT1",
    "EXT2", "EXT3", "EXT4", "EXT5", "EXT6", "EXT7", "EXT8",
    "EXT9", "EXT10"), agreeability = c("AGR1", "AGR2", "AGR3",
    "AGR4", "AGR5", "AGR6", "AGR7", "AGR8", "AGR9", "AGR10"),
  natural_reactions = c("EST1", "EST2", "EST3", "EST4", "EST5",
    "EST6", "EST7", "EST8", "EST9", "EST10"))

```

We see mean and standard deviation of each trait's questions. This shows us significant differences, as shown during EDA, and will be useful to use normality for estimating probabilities or score (percentiles) as defined by the standard test.

```

traitsMeans <- c(mean(rowMeans(df[, 2:11])), mean(rowMeans(df[,
  12:21])), mean(rowMeans(df[, 22:31])), mean(rowMeans(df[,
  32:41])), mean(rowMeans(df[, 42:51])))
traitsSds <- c(sd(rowMeans(df[, 2:11])), sd(rowMeans(df[, 12:21])),
  sd(rowMeans(df[, 22:31])), sd(rowMeans(df[, 32:41])), sd(rowMeans(df[,
  42:51])))
names(traitsMeans) <- c("extroversion_EXT", "natural_reactions_EST",
  "agreeability_AGR", "conscienciousness_CSN", "openess_OPN")
names(traitsSds) <- c("extroversion_EXT", "natural_reactions_EST",
  "agreeability_AGR", "conscienciousness_CSN", "openess_OPN")
traitsMeans %>%
  knitr::kable(digits = 4)

```

	x
extroversion_EXT	2.9595
natural_reactions_EST	3.0666
agreeability_AGR	3.7736
conscienciousness_CSN	3.3661
openess_OPN	3.8989

```

traitsSds %>%
  knitr::kable(digits = 4)

```

	x
extroversion_EXT	0.9104
natural_reactions_EST	0.8586
agreeability_AGR	0.7293
conscientiousness_CSN	0.7373
openess_OPN	0.6302

Dataset size reduction

Although Exploratory Data Analysis was done using all data available, we are now going to randomly reduce data size. This is done for computability and is defined in a single parameter *nObsersDevelopment* that can be increased at anytime if system capacity allows.

```
nObsersDevelopment <- 5000
set.seed(1, sample.kind = "Rounding")
df <- df[sample(nrow(df), nObsersDevelopment), ]
```

Save input data for shiny app

The shiny app is an interactive implementation of this model. It simplifies some parts of the process and shows prediction generated for each questions, as well as the calculated scores. Its purpose is educational and as a demo, it does not store data nor submit specific transaction, i.e. allows simulation with different values.

For this version of the web app we pass reduced version of the dataset to allow online training and prediction. There are significant improvements that could be added to the tool, which are remarked in the *Conclusions* chapter of this document.

```
BFdata <- as(data.matrix(df[, 2:51]), "realRatingMatrix")
saveRDS(BFdata, "./data_source/BFdata.rds")
saveRDS(dictionary, "./data_source/dictionary.rds")
```

Separate data in partitions

After data exploration, cleanup, adjustment, and reduction, we will separate in two partitions so we can separately train and test our models. We have the option to use *built-in* partition, testing and accuracy functions of the package recommenderlab, but for clarity and control, we will manage directly and independently these data partitions.

```
devReduction <- 0.2 # Percentage of original data we extract for development
set.seed(1, sample.kind = "Rounding")
test_index <- sample(c(1:nrow(df)), (nrow(df) * devReduction),
  replace = FALSE)
BF_train <- df[-test_index, ]
BF_test <- df[test_index, ]
totRowsValidation <- nrow(BF_test) # number of observations to predict
rm(test_index)
```

3.3 Training and prediction

Accuracy measurement: metrics of prediction success

A good definition of the accuracy we expect from our model is a key step during development. The multi-result (5 Scores) characteristic of this test outcomes makes it necessary to define our expectation for the predictions in terms of:

- 1) What exactly we want to measure for accuracy
- 2) A reference of how good the prediction is when done randomly

What to measure? The challenge of the project is to predict Scores for each one of the five personality traits, but these scores directly depend on the answers to 50 questions (10 per trait) where we only know real value of five of them. Thus, our accuracy will depend on how close we are to the real Score for each trait, which is originally expressed in percentile. Since this measure can be understood as approximate/soft, we assume that, for a specific trait, hitting (predicting) the **correct quartile** could be considered as a *reasonable good result*. In similar terms, hitting **same half (high-low)** is *also -though less- a good result*. Looking at all five trait as a set, we will also establish as a *good result* to correctly predict **(high-low) most traits**, i.e. three or more of the five. So, our accuracy metrics will be defined as:

Accuracy Indicator	Definition
‘Hits quartile’ for a trait	Score for a specific trait predicts correct quartile (1-25, 26-50, 51-75, 76-100)
‘Hits quartile’ all traits	average of <i>Hits quartile</i> for all traits of an observation (test)
‘Hits HighLow’ for a trait	Score for a specific trait predicts correct half (1-50, 51-100)
‘Hits HighLow all traits’	average of <i>Hits quartile</i> for all traits of an observation (test)
‘3+ hits HighLow’	Scores of three or more traits predict correctly halves

How good would be a random estimate? (Montecarlo approach) As reference for accuracy improvement, we will estimate what the Scores would be if using just a random criteria for predicting answers to each question. The random distribution of answers to 50 questions with 5 possible answers same probability (1/5) follow a binomial (Bernoulli) pattern. However, due to the calculations required for our accuracy indicators, we will code a **Montecarlo** simulation of a random selection, which will allow us to simulate results. These accuracies will be considered then as “base reference” for further improvements during modeling process.

This function adjusts random score taking into account that answers of one of every 10 question is known. For simplicity, we assume a linear improvement of accuracy for 1/10 and correct by chance of randomly hitting (1/5)

```
adjustScore <- function(x) {  
  x + (4/5) * ((realScores - x)/10)  
}  
# Generate a real result for our simulation  
set.seed(1, sample.kind = "Rounding")  
realScores <- sample(c(0:100), 5, replace = TRUE)  
realQuartiles <- 1 + floor(realScores/25)  
realHalves <- 1 + floor(realScores/50)  
# iterate simulated results obtained randomly  
B <- 10000  
estimAccuracy <- NULL  
for (b in c(1:B)) {
```

```

predictedScores <- sample(c(0:100), 5, replace = TRUE)
predictedScores <- round(adjustScore(predictedScores), 0)
predictedQuarters <- 1 + floor(predictedScores/25)
predictedHalves <- 1 + floor(predictedScores/50)
hitsQuartile <- (realQuartiles == predictedQuarters)
hitsHiLo <- (realHalves == predictedHalves)
mostHitHiLo <- (sum(hitsHiLo) >= 3)
estimAccuracy <- rbind(estimAccuracy, c(hitsQuartile, hitsHiLo,
  mostHitHiLo))
}
colnames(estimAccuracy) <- c(c("hitsQuartile_0", "hitsQuartile_C",
  "hitsQuartile_E", "hitsQuartile_A", "hitsQuartile_N"), c("hitsHiLo_0",
  "hitsHiLo_C", "hitsHiLo_E", "hitsHiLo_A", "hitsHiLo_N"),
  "MostHitHiLo")
estimAccuracy <- colMeans(estimAccuracy)
# We just obtained accuracies for each category Create a
# dataframe to store results of the analysis
analysis_results <- data_frame(Trait = "All", Score = estimAccuracy["MostHitHiLo"],
  Accuracy_type = "3+ hits HighLow", Algorithm = "Montecarlo") %>%
  rbind(data_frame(Trait = "O_score", Score = estimAccuracy["hitsQuartile_0"],
    Accuracy_type = "Hits quartile", Algorithm = "Montecarlo"),
    data_frame(Trait = "C_score", Score = estimAccuracy["hitsQuartile_C"],
      Accuracy_type = "Hits quartile", Algorithm = "Montecarlo"),
    data_frame(Trait = "E_score", Score = estimAccuracy["hitsQuartile_E"],
      Accuracy_type = "Hits quartile", Algorithm = "Montecarlo"),
    data_frame(Trait = "A_score", Score = estimAccuracy["hitsQuartile_A"],
      Accuracy_type = "Hits quartile", Algorithm = "Montecarlo"),
    data_frame(Trait = "N_score", Score = estimAccuracy["hitsQuartile_N"],
      Accuracy_type = "Hits quartile", Algorithm = "Montecarlo"),
    data_frame(Trait = "O_score", Score = estimAccuracy["hitsHiLo_0"],
      Accuracy_type = "Hits HighLow", Algorithm = "Montecarlo"),
    data_frame(Trait = "C_score", Score = estimAccuracy["hitsHiLo_C"],
      Accuracy_type = "Hits HighLow", Algorithm = "Montecarlo"),
    data_frame(Trait = "E_score", Score = estimAccuracy["hitsHiLo_E"],
      Accuracy_type = "Hits HighLow", Algorithm = "Montecarlo"),
    data_frame(Trait = "A_score", Score = estimAccuracy["hitsHiLo_A"],
      Accuracy_type = "Hits HighLow", Algorithm = "Montecarlo"),
    data_frame(Trait = "N_score", Score = estimAccuracy["hitsHiLo_N"],
      Accuracy_type = "Hits HighLow", Algorithm = "Montecarlo"),
    data_frame(Trait = "All", Score = mean(estimAccuracy[1:5]),
      Accuracy_type = "Hits quartile", Algorithm = "Montecarlo"),
    data_frame(Trait = "All", Score = mean(estimAccuracy[6:10]),
      Accuracy_type = "Hits HighLow", Algorithm = "Montecarlo"))
# Show results (theoretical random estimate with a
# montecarlo approach) as base reference for next
# improvements during modelling
analysis_results %>%
  knitr::kable(digits = 4)

```

Trait	Score	Accuracy_type	Algorithm
All	0.5398	3+ hits HighLow	Montecarlo
O_score	0.2665	Hits quartile	Montecarlo
C_score	0.2684	Hits quartile	Montecarlo

Trait	Score	Accuracy_type	Algorithm
E_score	0.2744	Hits quartile	Montecarlo
A_score	0.2680	Hits quartile	Montecarlo
N_score	0.2499	Hits quartile	Montecarlo
O_score	0.5127	Hits HighLow	Montecarlo
C_score	0.4999	Hits HighLow	Montecarlo
E_score	0.5167	Hits HighLow	Montecarlo
A_score	0.5434	Hits HighLow	Montecarlo
N_score	0.5277	Hits HighLow	Montecarlo
All	0.2654	Hits quartile	Montecarlo
All	0.5201	Hits HighLow	Montecarlo

Questions selection: Minimum Combined Correlation Algorithm (MCCA)

Our model relies on **only five questions** actually known, out of 50, to predict the other 45. This means a good selection of the five questions we get real answer for, is very important to better predict the others. However, choosing just the same apparent best combination (always the same for given dataset) seems not to be the most open, realistic, rich option, since in case of moving this model into production, would result in a very poor diversity of input results, which long term would drive to worse results and lack of diversity. Thus, we will state as premise that, one of the questions to get real answer from, must be randomly selected (seed question), whilst the other four ones can be generated based on expected best results.

The algorithm we propose to select those other four question can be called “minimum combined correlation method”, and consists in selecting those questions having less (absolute) correlation with the seed question and with the other chose questions. It relies on the assumption that “the most *different* our know questions are, the better prediction we eill generate”We implement therefore a recursive approach based in comparing the average correlation of all pairs of questions for each combination (10) of potential questions including the seed question (10000)

```
# Minimum combined correlation algorithm: Generate all
# possible combinations of questions (10^5)
potentialQuestionsSets <- as(expand.grid(keys.list.allPositive),
  "matrix")
# generate first 'seed' question
set.seed(1, sample.kind = "Rounding")
randomSeedQuestion <- colnames(BF_test[, 2:51])[sample(1:50,
  1)]
randomSeedQuestion
```

```
## [1] "EST4"
```

```
# Filter only sets of questions containing the initial
# ramdom seed question
potentialQuestionsSets <- potentialQuestionsSets[which(rowAnys(potentialQuestionsSets ==
  randomSeedQuestion)), ]
# Once we know potential combinations of questions to get
# answer from, we need to calculate which is best for
# predicting. Our premise will be to get the set with
# minimum correlation among its pairs. The reason is that
# low correlation will inform us better of the 'difficult'
# questions where model will be weaker
correlPerSet <- NULL
```

```

# Loop each potential set of questions
for (qs in 1:nrow(potentialQuestionsSets)) {
  # generate all pairs of questions (10) within given set
  couplesThisCombination <- combinations(5, 2, v = potentialQuestionsSets[qs,
  ])
  pairsCorrelationsThisSet <- NULL
  # now we loop those pairs to get a summary number based
  # on their correlation
  for (c in 1:nrow(couplesThisCombination)) {
    question1 <- couplesThisCombination[c, 1]
    question2 <- couplesThisCombination[c, 2]
    pairsCorrelationsThisSet <- c(pairsCorrelationsThisSet,
      abs(corAllQuestions[question1, question2]))
  }
  correlPerSet <- c(correlPerSet, mean(pairsCorrelationsThisSet))
}

```

Based on the minimum correlation among its questions, we select a set:

```

# Convert it to column index for continuing modelling
# process
chosen_questions <- which(colnames(BF_test[, 2:51]) %in% potentialQuestionsSets[which.min(correlPerSet)
  ])
chosen_questions

```

```
## [1] 8 14 24 36 41
```

Loop available recommendation algorithms

Our approach for predicting 45 answers (items) of an user that actually answered only to 5, will be to use a Recommendation Model, based in the **library Recommenderlab**. The basic idea is to consider Users (those who answer) and Items (questions) and apply some of the methods of standard recommendation. However, there are some singularities in our case compared to a *typical movie/book* recommendation case:

- No sparcity: our matrix of answers *Users X Items* has all cells filled, since we have data for all users answering all questions.
- Matrix is pretty “vertical” having more that 800000 Users (after initial cleanup) and 50 Item (columns).
- Items have clear correlated components, since questions are grouped by traits.

Therefore, we will choose adequate algorithms based on singularities:

Due to the shape of the matrix, we discard IBCF (Item Based collaborative Filtering) because of the risk in some cases for not generating predictions for all item (could be fixed with average/median filling though loosing accuracy). Regarding no scarcity we also discard Popular method.

In order to generate a rich approach for our prediction, we will focus in a collaborative filtering method (UBCF) and in a matrix factorization method (ALS) - we choose it instead of SVD for its superior accuracy given singularities of our User x Item matrix:

UBCF (User Based Collaborative Filtering): this algorithm mimics word-of-mouth by analyzing rating data from many individuals. The assumption is that users with similar preferences will rate (answer) items (questions) similarly. Thus missing ratings for a user can be predicted by first finding a *neighborhood* of similar users and then aggregate the ratings of these users to form a prediction.

The neighborhood is defined in terms of similarity between users. Popular similarity measures for are the *Pearson correlation coefficient* and the *Cosine similarity* (the one we will use). This similarity measure is defined between two users u_x and u_y as:

$$sim_{\cos}(\vec{x}, \vec{y}) = \frac{\vec{x} \cdot \vec{y}}{\|\vec{x}\| \cdot \|\vec{y}\|},$$

where $\vec{x} = r_x$ and $\vec{y} = r_y$ represent the row vectors in R with the two users' profile vectors. $sd(\hat{u})$ is the standard deviation and $\|\cdot\|$ is the l^2 -norm of a vector. For calculating similarity using rating data only the dimensions (items) are used which were rated by both users.

Now the neighborhood for the active user $\mathcal{N}(a) \subset \mathcal{U}$ can be selected by either a threshold on the similarity or by taking the k nearest neighbors. Once the users in the neighborhood are found, their ratings are aggregated to form the predicted rating for the active user. The easiest form is to just average the ratings in the neighborhood.

$$\hat{r}_{aj} = \frac{1}{|N(a)|} \sum_{i \in N(a)} r_{ij}$$

ALS (Alternating Least Squares): ALS is an iterative optimization process where, for every iteration, we try to arrive closer and closer to a factorized representation of our original data. We have our original matrix R of size $u \times i$ with our users, items and some type of feedback data. We then want to find a way to turn that into one matrix with users and hidden features of size $u \times f$ and one with items and hidden features of size $f \times i$. In U and V we have weights for how each user/item relates to each feature. What we do is we calculate U and V so that their product approximates R as closely as possible: $R \approx U \times V$. By randomly assigning the values in U and V and using least squares iteratively we can arrive at what weights yield the best approximation of R . The least squares approach in it's basic forms means fitting some line to the data, measuring the sum of squared distances from all points to the line and trying to get an optimal fit by minimising this value.

With the alternating least squares approach we use this idea but iteratively alternate between optimizing U and fixing V and vice versa. We do this for each iteration to arrive closer to $R = U \times V$.

```
# we prepare methods to run
methods_choice <- list(list("ALS", ALS = list(NULL)), list("UBCF",
  'user-based CF' = list(nn = 50)))

# This for-loop serves for exution of each chosen method to
# train, test and measure-accuracy
for (a in 1:length(methods_choice)) {
  # Train model (Recommenderlab package)
  chosenAlgorithm <- methods_choice[[a]][[1]]
  chosenAlgorithmParams <- methods_choice[[a]][[2]]
  # Train
  recom <- Recommender(as(data.matrix(BF_train[, 2:51]), "realRatingMatrix"),
    method = chosenAlgorithm, parameter = chosenAlgorithmParams)

  # Test (validate) model
  questionsList <- dictionary$Question
  names(questionsList) <- dictionary$ID

  # Prepare known ratings (5 out of 50) to send to the
  # model - take from VALIDATION
  ratings <- matrix(NA, nrow = totRowsValidation, ncol = 50)
  ratings[, chosen_questions[1]] <- BF_test[, 2:51][, chosen_questions[1]]
}
```



```

ratings[, chosen_questions[2]] <- BF_test[, 2:51][, chosen_questions[2]]
ratings[, chosen_questions[3]] <- BF_test[, 2:51][, chosen_questions[3]]
ratings[, chosen_questions[4]] <- BF_test[, 2:51][, chosen_questions[4]]
ratings[, chosen_questions[5]] <- BF_test[, 2:51][, chosen_questions[5]]
ratings <- as(ratings, "realRatingMatrix")
# Predict create (predict) recommendations (45) based
# on known 5 answers ('ratings')
pred <- predict(recom, ratings, n = 45)
# Predicted answers from the model (still 'rough')
matrixNamesPredicted <- matrix(unlist(getList(pred)), ncol = 45,
                               byrow = TRUE)
matrixScoresPredicted <- matrix(unlist(getRatings(pred)),
                                ncol = 45, byrow = TRUE)
# This matrix (matrixScoresPredicted) contain all
# predictions for validation dataset row by row
# (matrixScoresPredicted), BUT each row follows a
# different order, based on corresponding row of
# previous (matrixNamesPredicted) matrix So, we need to
# rearrange all lines to any, but the same, column
# structure. We use list of first row as template for
# this arrangement:
tmpPatternColumnsReference <- matrixNamesPredicted[1, ]
tmpPatternColumnsReference
# This loop performs rearrange row by row
for (r in c(1:totRowsValidation)) {
  matrixScoresPredicted[r, ] <- matrixScoresPredicted[r,
    ][match(tmpPatternColumnsReference, matrixNamesPredicted[r,
    ])]
}
# We assign column names of first row (used as
# reference to new global matrix)
colnames(matrixScoresPredicted) <- tmpPatternColumnsReference
rm(matrixNamesPredicted) # Remove this mixed matrix to prevent confusion,
# clarity in next steps Inspect obtained matrix with
# sorted (align values per column) predictions
head(matrixScoresPredicted)

# Join together predictions with 5 known answers Let's
# collect together entered answers and predicted
# ratings to prepare results calculation
realRatings <- BF_test[, 2:51] # real n x 50 answers in the validation dataset
head(realRatings)
# names of 5 'known':
enteredQuestions <- dictionary[chosen_questions[1:5], 1]
# answers of n x 5 'known' questions:
enteredRatings <- as(ratings, "matrix")[, chosen_questions[1:5]]
colnames(enteredRatings) <- enteredQuestions
# Build matrix as the union of 5 real answers + 45
# predicted for all validation rows
dim(enteredRatings) # n x 5 questions entered
dim(matrixScoresPredicted) # n x 45 questions predicted
# Ready to prepare a single matrix
matrixAllRatings <- cbind(enteredRatings, matrixScoresPredicted)

```

```

dim(matrixAllRatings) # n x 50 (all questions)
rm(matrixScoresPredicted) # For clarity to prevent confusion
colnames(matrixAllRatings)
# Reorder 'mixed' columns as in dataset
matrixAllRatings <- matrixAllRatings[, colnames(BFdata)]
colnames(matrixAllRatings)

# Score real and predicted matrices We have now both
# matrices (real data and predicted data) so we score
# them individually Score adding last (predicted) row
# to validation data
scoresfile_real <- scoreFast(keys.list.allPositive, as(BF_test[,
  2:51], "matrix"))
scoresfile_pred <- scoreFast(keys.list.allPositive, as(matrixAllRatings,
  "matrix"))
# Put a suffix in predicted-based scores columns to
# distinguish from real-based
colnames(scoresfile_pred) <- paste(colnames(scoresfile_real),
  "_pred", sep = "")

# Convert just calculated scores to ranking
# (percentile) Results are presented in percentile for
# each user/trait. With all data + prediction scored
# together (n+1 rows matrix), calculate percentiles of
# the last row (results for the user) using eCDF For
# real data percentiles
O_score <- round(ecdf(scoresfile_real[, "openness-A"])(scoresfile_real[,
  "openness-A"]) * 100, 0)
C_score <- round(ecdf(scoresfile_real[, "conscientiousness-A"])(scoresfile_real[,
  "conscientiousness-A"]) * 100, 0)
E_score <- round(ecdf(scoresfile_real[, "extroversion-A"])(scoresfile_real[,
  "extroversion-A"]) * 100, 0)
A_score <- round(ecdf(scoresfile_real[, "agreeability-A"])(scoresfile_real[,
  "agreeability-A"]) * 100, 0)
N_score <- round(ecdf(scoresfile_real[, "natural_reactions-A"])(scoresfile_real[,
  "natural_reactions-A"]) * 100, 0)
traits_percentiles_real <- cbind(O_score, C_score, E_score,
  A_score, N_score)
# For predicted data percentiles
O_score_pred <- round(ecdf(scoresfile_real[, "openness-A"])(scoresfile_pred[,
  "openness-A_pred"]) * 100, 0)
C_score_pred <- round(ecdf(scoresfile_real[, "conscientiousness-A"])(scoresfile_pred[,
  "conscientiousness-A_pred"]) * 100, 0)
E_score_pred <- round(ecdf(scoresfile_real[, "extroversion-A"])(scoresfile_pred[,
  "extroversion-A_pred"]) * 100, 0)
A_score_pred <- round(ecdf(scoresfile_real[, "agreeability-A"])(scoresfile_pred[,
  "agreeability-A_pred"]) * 100, 0)
N_score_pred <- round(ecdf(scoresfile_real[, "natural_reactions-A"])(scoresfile_pred[,
  "natural_reactions-A_pred"]) * 100, 0)
traits_percentiles_predicted <- cbind(O_score_pred, C_score_pred,
  E_score_pred, A_score_pred, N_score_pred)

# Here we have all resulting data

```

```

names(chosenAlgorithm) = "Algorithm"
# Other files generated during process
head(traits_percentiles_real)
head(traits_percentiles_predicted)

# Compose high level result in a data frame We need to
# calculate several accuracy indicators (by
# trait/general, by quarter/half correct prediction,
# and by most traits correctly predicted) Detect
# quartile success in trait
accuracyPerQuartile <- colMeans((1 + floor(abs((traits_percentiles_real -
  1))/25)) == (1 + floor(abs((traits_percentiles_predicted -
  1))/25)))
accuracyPerQuartileMean <- as(cbind("All", mean(accuracyPerQuartile),
  "Hits quartile"), "matrix")
accuracyPerQuartile <- cbind(as(accuracyPerQuartile, "matrix"),
  Accuracy_type = "Hits quartile")
# Detect High-Low success in trait
accuracyPerHalf <- colMeans((1 + floor(abs((traits_percentiles_real -
  1))/50)) == (1 + floor(abs((traits_percentiles_predicted -
  1))/50)))
accuracyPerHalfMean <- as(cbind("All", mean(accuracyPerHalf),
  "Hits HighLow"), "matrix")
accuracyPerHalf <- cbind(as(accuracyPerHalf, "matrix"), Accuracy_type = "Hits HighLow")
# Detect High-Low successes are >= 3 in same test
accuracySameHalf <- (1 + floor(abs((traits_percentiles_real -
  1))/50)) == (1 + floor(abs((traits_percentiles_predicted -
  1))/50))
accuracySameHalf <- mean(rowSums(accuracySameHalf[, 1:5]) >=
  3)
accuracySameHalf <- as(cbind("All", accuracySameHalf, "3+ hits HighLow"),
  "matrix")
accuracySameHalf

# Collect results already obtained in a temporary
# matrix
tmpResultsPrediction <- rbind(accuracyPerQuartile, accuracyPerHalf)
tmpResultsPrediction <- cbind(Trait = rownames(tmpResultsPrediction),
  tmpResultsPrediction)
colnames(tmpResultsPrediction)[2] <- "Score"
rownames(tmpResultsPrediction) <- NULL
tmpResultsPrediction <- tmpResultsPrediction %>%
  rbind(accuracySameHalf, accuracyPerQuartileMean, accuracyPerHalfMean) %>%
  cbind(Algorithm = chosenAlgorithm)
head(tmpResultsPrediction)
# Create data frame
df_resultsPrediction <- as.data.frame(tmpResultsPrediction)
df_resultsPrediction <- df_resultsPrediction %>%
  mutate(Score = as.double(as.character(df_resultsPrediction$Score)))
head(df_resultsPrediction)
rm(tmpResultsPrediction)
# Store results
analysis_results <- bind_rows(analysis_results, df_resultsPrediction)

```

```
analysis_results %>%  
  knitr::kable(digits = 4)  
}  
# Ends FOR-LOOP of algorithms to model
```

4 The shiny application

In addition to the model created based on existing dataset, we build also an interactive tool to show in practice how prediction works. It is fully developed in R as part of current project and provisioned to internet through platform Shinyapps and can be publicly accessed at <https://juanesuarez.shinyapps.io/BigFive5q/>

This tool is basically an interactive web to allow any person to answer 5 randomly selected questions, and based on them, estimate personality traits Big Five Scores by *predicting* all other 45 remaining question using our recommendation model.




Figure 2: Screenshot1 of interactive tool

Although this tool is fully operational, for practical and academic reasons, it is simplified to allow wide interactive testing avoiding any personal data protection issues. Some of the simplification changes:

- An user can repeat test as many times as desired.
- Algorithm implemented is UBCF (second best in our study) instead of ALS for interactivity performance.
- Questions are chosen randomly instead of using our just developed *Minimum Combined Correlation Algorithm*.
- All *hidden* predicted question's answers are shown for instructional purposes.

This promising tool inspires many developments to incorporate nice functionalities, potentially very useful in combination with other tools for other specific needs, Some of them:

- Allow multilanguage.
- Add a submission process to manage transactionally users and their results, preventing repetition and storing results.
- Hide results of prediction of unknown questions.
- Combine with other matter-related *control/classification" questions to improve prediction or link to other new outcomes.

5 Results

We see how accuracies, as defined, get significantly improved with both ALS and UBCF methods algorithms respect to base “random” reference generated with the Montecarlo approach.

We have got a very appreciable result after training and validating our two chosen algorithms (UBCF and ALS). UBCF offers better performance but ALS results in improves accuracies for all indicators.

Accuracy_type	Algorithm	All	O_score	C_score	E_score	A_score	N_score
3+ hits HighLow	ALS	0.874					
3+ hits HighLow	Montecarlo	0.5398					
3+ hits HighLow	UBCF	0.841					
Hits HighLow	ALS	0.7194	0.708	0.744	0.758	0.749	0.638
Hits HighLow	Montecarlo	0.52008	0.5127	0.4999	0.5167	0.5434	0.5277
Hits HighLow	UBCF	0.697	0.693	0.717	0.745	0.7	0.63
Hits quartile	ALS	0.4422	0.411	0.45	0.486	0.496	0.368
Hits quartile	Montecarlo	0.26544	0.2665	0.2684	0.2744	0.268	0.2499
Hits quartile	UBCF	0.4152	0.429	0.417	0.432	0.431	0.367

Summarizing by type of accuracy, results show clear improvements respect to base reference:

Accuracy_type	ALS	Montecarlo	UBCF	Improvement_UBCF	Improvement_ALS
3+ hits HighLow	0.874	0.5398	0.841	156%	162%
Hits HighLow	0.7194	0.52008	0.697	134%	138%
Hits quartile	0.4422	0.26544	0.4152	156%	167%

As a final high level number of prediction improvement, the improvement in global accuracy **reach an impressive 162% when all measured accuracies consolidated.**

6 Conclusion

The model resulting of this project is an inspiring method to dramatically reduce, with relatively low impact on results, the cost of filling very long questionnaires to assess personality traits.

We have followed unchanged the definition of a widely used and prestigious “Big Five Traits Test”, but adding data science to improve the process. More specifically, we have used thousands of existing test to predict answers of future users in an absolutely simplified way.

Probably the most exciting outcome of this project is that all this modeling is easy to integrate to solve problems in many areas where personality of individuals *play a role* in the problem description or prediction. Also, this work open the possibility to adapt similar approach to other psychometric challenges.

As an *extra bonus*, we have build also a web app to interactively show results and make it visible how model works. This tool wants to catalize inspiration of data scientist and general viewers on the many possibilities we have in front of us.

Measurement of the accuracy of the model has required to define and calculate indicators of how good the model works in predicting the right trait score for an individual. The fact result is a set of five numbers “percentiles” will make us define metrics for assessing how close predictions are to real results.

There are some clear **premises we strictly followed for this project**:

- Measurable: Strict measurement of results, based on a clear definition of the accuracy obtained.
- Useful and escalable: Prepare model to be useful for further developments/integrations, i.e. make it as a *pluggable module* to be integrated in other systems. For example, instead of showing always “optimal” questions to be answered, we prefer to make them -partially- random, so a deployment at bigger scale might benefit of variability.
- Independent validation: Never use validation (test) data partition to estimate parameters of select any criteria during training process
- Computability: project is deep in calculation but affordable within a local environment in terms of performance. This means that we reduced our dataset during some stages, but a wider execution can be always an option without need of changing source code, just the parameter.

There are several **original aspects** on this project:

- Use Recommender: To take advantage of a standard “recommendation model”, adapting some concepts to use classical “User x Item” as inspiration for a “User x Question/Trait” structure. The advantage is to use strong prediction tool (i.e. Recommenderlab, etc.), but on the other side fact like the lack of matrix sparcicity or the very “vertical” shape of the ratings matrix (approx. 10^6 users x only 50 columns) require specific attention during the model implementation.
- An app to play: To create an interactive, open access, graphical app tool to facilitate understanding of the concept and even disclosing some of the key prediction steps followed.
- Questions selection: Our own algorithm to optimize selection of questions to get answer for, so the change (to ensure varibility on data retrieval), whilst optizing combination so final accuracy is best.

Furthermore, the promising methods created and results obtained open the door to potential developments/additions for future versions:

- Keep optimizing/tuning Recommender (evaluationScheme, k-fold, bootstrap, no-sparse specifics, ...).
- Increase dataset size according to computation capability available.
- Ensemble to use different algorithm optimal for each specific trait score prediction.
- Ensemble a model pre or post recommender to consider Month and/or Country as significant predictors.
- Extend to 100 questions version of the Test and/or increase 5 to more questions to use as base.

- Add external questions to facilitate classification.
- Integrate model with existing model that already analyzed outcomes predicted by personality traits.
- Move web app to “production” doing improvements and allowing public users functionality.

As a final conclusion, we can state we have successfully accomplished challenge goal in terms of project requirements and accuracy for the validation set provided.